



NATURAL LANGUAGE PROCESSING

Arpendu Ganguly

NATURAL LANGUAGE PROCESSING

- **Natural language processing (NLP) is any computation, manipulation of natural language**
- Natural Language
 - Language used for everyday communication by humans
 - Evolves with time
- **Text data is growing fast!**
 - Data continues to grow exponentially
 - Approximately 80% of all data is estimated to be unstructured, text-rich data

NATURAL LANGUAGE PROCESSING

Parse text

Find / Identify / Extract
relevant information from text

Classify text documents

Search for relevant text
documents

Sentiment analysis

Topic modeling

Data hidden in plain sight

The image shows a screenshot of the Twitter profile for the UN Spokesperson (@UN_Spokesperson). The profile header includes the name, handle, bio, location (New York, USA), and join date (May 2010). The bio states: "Official Twitter account of the Office of the Spokesperson for United Nations Secretary-General Ban Ki-moon." The location is "New York, USA" and the join date is "Joined May 2010". The profile statistics show 14.6K tweets, 994 following, 391K followers, 49 likes, and 3 lists. The profile picture is the UN logo. The header image shows a man speaking at a podium. The tweets section shows three tweets. The first tweet is about maintaining unity on the Korean Peninsula. The second tweet is about ethics being built into the ideals and objectives of the United Nations. The third tweet is about the UN family being fortunate to have a wonderful supporter. Annotations with green boxes and arrows point to various parts of the profile and tweets, highlighting data hidden in plain sight.

Annotations:

- Social network:** Points to the Twitter profile header.
- Author:** Points to the profile picture (UN logo).
- Description:** Points to the bio text.
- Location:** Points to the location text.
- Tweet:** Points to the tweet text, with sub-annotations for:
 - Topic:** Points to the main text of the tweet.
 - Sentiment:** Points to the tweet text.
 - Time:** Points to the timestamp of the tweet.
 - Popularity:** Points to the engagement metrics (retweets, replies, likes).

NPL TASKS

NLP Tasks: A Broad Spectrum

- Counting words, counting frequency of words
- Finding sentence boundaries
- Part of speech tagging
- Parsing the sentence structure
- Text Classification, Identifying semantic roles
- Identifying entities in a sentence / Named Entity Recognition
- Finding which pronoun refers to which entity / Co-reference and pronoun resolution
- much more ...

NLTK: Natural Language Toolkit

- Open source library in Python
- Has support for most NLP tasks
- Also provides access to numerous text corpora

TOKENIZATION

Splitting a sentence into words
/ tokens

```
test1 = "Children shouldn't drink a sugary drink before bed."
```

```
test1.split(" ")
```

```
['Children', "shouldn't", 'drink', 'a', 'sugary', 'drink', 'before', 'bed.']
```

```
len(test1.split(" "))
```

8

```
nlk.word_tokenize(test1)
```

```
['Children',  
'should',  
"n't",  
'drink',  
'a',  
'sugary',  
'drink',  
'before',  
'bed',  
'']
```

```
len(nltk.word_tokenize(test1))
```

10

COUNTING VOCABULARY OF WORDS

Count unique words

```
text="Children shouldn't drink a sugary drink before bed."
```

```
len(text)
```

```
51
```

```
nltk.word_tokenize(text)
```

```
['Children',  
 'should',  
 "n't",  
 'drink',  
 'a',  
 'sugary',  
 'drink',  
 'before',  
 'bed',  
 '.']
```

```
len(nltk.word_tokenize(text))
```

```
10
```

```
len(set(nltk.word_tokenize(text)))
```

```
9
```

```
list(set(nltk.word_tokenize(text))[:4])
```

```
['before', '.', 'sugary', 'bed']
```

FREQUENCY OF WORDS

Mapping of unique word to count

```
text1="I felt happy because I saw the others were happy and because I knew I should feel happy, but I wasn't really happy."
```

```
dist = FreqDist(nltk.word_tokenize(text1))  
dist
```

```
FreqDist({' ': 1,  
          '.': 1,  
          'I': 5,  
          'and': 1,  
          'because': 2,  
          'but': 1,  
          'feel': 1,  
          'felt': 1,  
          'happy': 4,  
          'knew': 1,  
          'others': 1,  
          'really': 1,  
          'saw': 1,  
          'should': 1,  
          't': 1,  
          'the': 1,  
          'wasn': 1,  
          'were': 1,  
          ''': 1}))
```

FREQUENCY OF WORDS

Mapping of unique word to count

```
dist.keys()
```

```
dict_keys(['I', 'felt', 'happy', 'because', 'saw', 'the', 'others', 'were', 'and', 'knew', 'should', 'feel', ',', 'but', 'was', 'n', "'", 't', 'really', '.'])
```

```
dist.values()
```

```
dict_values([5, 1, 4, 2, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1])
```

```
vocab1 = dist.keys()  
list(vocab1)[:4]
```

```
['I', 'felt', 'happy', 'because']
```

```
dist["I"]
```

```
5
```

```
for w in vocab1:  
    if len(w) > 3 and dist[w] > 3:  
        print(w)
```

```
happy
```

```
freqwords = [w for w in vocab1 if len(w) > 3 and dist[w] > 3]  
freqwords
```

```
['happy']
```


NORMALIZATION AND STEMMING

Normalization involves eliminating punctuation, converting the entire text into lowercase or uppercase and so on.

```
input1 = "List listed lists listing listings"  
input1
```

```
'List listed lists listing listings'
```

```
words1=input1.lower().split(" ")  
words1
```

```
['list', 'listed', 'lists', 'listing', 'listings']
```

```
porter = nltk.PorterStemmer()  
  
#List Comprehension  
[porter.stem(t) for t in words1]
```

```
['list', 'list', 'list', 'list', 'list']
```

```
[porter.stem(t) for t in input1.split(" ")]
```

```
['list', 'list', 'list', 'list', 'list']
```

STEMMING

Reduce inflectional forms and sometimes derivationally related forms of a word to a common base form.

```
text=nltk.corpus.udhr.words('English-Latin1')[7:20]  
text
```

```
['recognition',  
'of',  
'the',  
'inherent',  
'dignity',  
'and',  
'of',  
'the',  
'equal',  
'and',  
'inalienable',  
'rights',  
'of']
```

```
[porter.stem(t) for t in text]
```

```
['recognit',  
'of',  
'the',  
'inher',  
'digniti',  
'and',  
'of',  
'the',  
'equal',  
'and',  
'inalien',  
'right',  
'of']
```

STEMMING AND LEMMATIZATION

Stemming and lemmatization

reduce inflectional forms and sometimes derivationally related forms of a word to a common base form.

Lemmatization: Stemming, but resulting stems are all valid words.

```
WNlemma = nltk.WordNetLemmatizer()  
[WNlemma.lemmatize(t) for t in text]
```

```
['Universal',  
'Declaration',  
'of',  
'Human',  
'Rights',  
'Preamble',  
'Whereas',  
'recognition',  
'of',  
'the',  
'inherent',  
'dignity',  
'and',  
'of',  
'the',  
'equal',  
'and',  
'inalienable',  
'right',  
'of']
```

PART-OF-SPEECH (POS) TAGGING

Part of speech tagging

- identification of words as nouns, verbs, adjectives, adverbs, etc
- Many more tags or word classes than just these

```
text11 = "Children shouldn't drink a sugary drink before bed."  
text13 = nltk.word_tokenize(text11)  
# NLTK's Tokenizer  
nltk.pos_tag(text13)
```

```
[('Children', 'NNP'),  
 ('shouldn't', 'MD'),  
 ('n't', 'RB'),  
 ('drink', 'VB'),  
 ('a', 'DT'),  
 ('sugary', 'JJ'),  
 ('drink', 'NN'),  
 ('before', 'IN'),  
 ('bed', 'NN'),  
 ('.', '.')] ]
```

```
: nltk.help.upenn_tagset('MD')
```

```
MD: modal auxiliary  
can cannot could couldn't dare may might must need ought shall should  
shouldn't will would
```

TEXT FEATURE EXTRACTION

Tf-idf transform/term weighting

- Often words occurring frequently (e.g. “the”, “a”, “is” in English) carry very little meaningful information about the actual contents of the document.
- Weights high to terms which are rarer yet more interesting
- **tf-idf** means **term-frequency times inverse document-frequency**
- **Tf** means **term-frequency**, the number of times a term occurs in a given document
- **Inverse document-frequency**

- where N is the total number of documents, and n_t is the number of documents that contain term t .

TEXT FEATURE EXTRACTION

Bag of Words representation

- **tokenizing** strings and giving an integer id for each possible token, for instance by using white-spaces and punctuation as token separators
- **counting** the occurrences of tokens in each document
- **normalizing and weighting** with diminishing importance tokens that occur in the majority of samples / documents

Sparse matrix

- sparse matrix or sparse array is a matrix in which most of the elements are zero

n -gram

- A contiguous sequence of n items from a given sample of text or speech
- Example
 - “I am working in Accenture.”
 - Unigram (1 gram): “I”, “am”, “working”, “in”, “Accenture”
 - Bigram (2 gram) : “I am”, “am working”, “working in”, “in Accenture”